

Contents

- [AA 203 HW 8 Question 3](#)
- [Part a - Plain LQR](#)
- [Part b- Certainty equivalent adaptive LQR controller](#)
- [Part c- adaptive LQR controller with white noise](#)

AA 203 HW 8 Question 3

Somrita Banerjee

```
clc
clear all
close all
rng('default') % For reproducibility

A0 = [0.99 0 0;
      0 0.99 0;
      0 0 0.99];
B0 = [1 0 0;
      0 1 0;
      0 0 1];
A = [1.01 0.01 0;
     0.01 1.01 0.01;
     0 0.01 1.01];
B = B0;

noiseStdDev = 0.01;
Q = eye(3);
R = 1000*eye(3);

xInit = [1.0; 1.5; 2.0];
```

Part a - Plain LQR

```
[K,S,e] = dlqr(A0,B0,Q,R);
N_vals = [100 1000 10000];
for N = N_vals
    x = zeros(3, N);
    u = zeros(3, N);
    cost = 0;
    for i = 1:N
        if i == 1
            x(:,i) = xInit;
        else
            noise = normrnd(0, noiseStdDev, [3,1]);
            x(:,i) = A*x(:,i-1) + B*u(:,i-1) + noise;
        end
        u(:,i) = -K*x(:,i);
        cost = cost + x(:,i)'*Q*x(:,i) + u(:,i)'*R*u(:,i);
    end

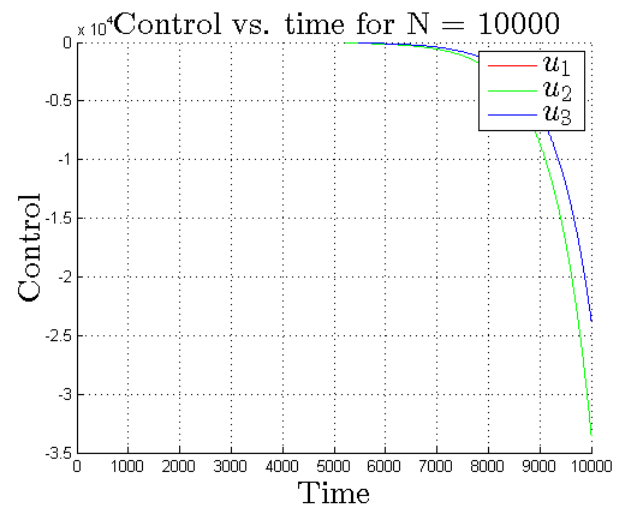
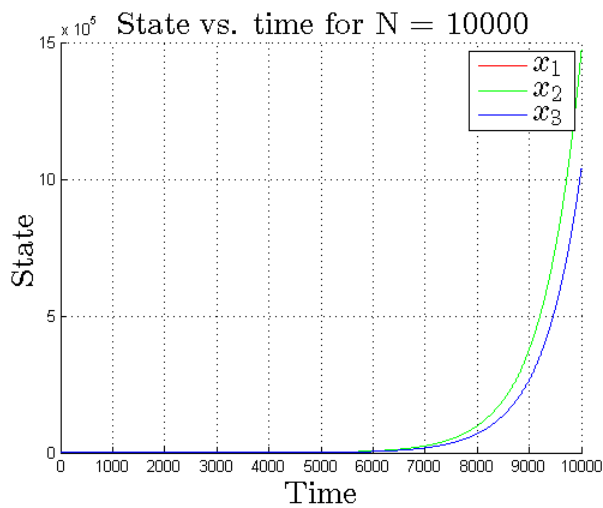




    fprintf('Cost for N = %d is %.2f \n',N, cost)

    figure
    hold on
    plot([1:N], x(1,:), 'r')
    plot([1:N], x(2,:), 'g')
    plot([1:N], x(3,:), 'b')
    titl = sprintf('State vs. time for N = %d',N);
    legend({'$x_1$', '$x_2$', '$x_3$'}, 'Interpreter', 'latex', 'FontSize', 20);
    xlabel('Time', 'Interpreter', 'latex', 'FontSize', 20);
    ylabel('State', 'Interpreter', 'latex', 'FontSize', 20);
    title(titl, 'Interpreter', 'latex', 'FontSize', 20);
    grid on

    figure
    hold on
    plot([1:N], u(1,:), 'r')
    plot([1:N], u(2,:), 'g')
    plot([1:N], u(3,:), 'b')
    titl = sprintf('Control vs. time for N = %d',N);
    legend({'$u_1$', '$u_2$', '$u_3$'}, 'Interpreter', 'latex', 'FontSize', 20);
    xlabel('Time', 'Interpreter', 'latex', 'FontSize', 20);
    ylabel('Control', 'Interpreter', 'latex', 'FontSize', 20);
    title(titl, 'Interpreter', 'latex', 'FontSize', 20);
    grid on
end
```

```
Cost for N = 100 is 1216.89
Cost for N = 1000 is 53255.71
Cost for N = 10000 is 2418241054959084.00
```

We see that our naive implementation of an LQR with the wrong parameters results in the cost increasing drastically as the number of time steps increases. The state and control also become unbounded.



Part b- Certainty equivalent adaptive LQR controller

```

N_vals = [100 1000 10000];
for N = N_vals
    x = zeros(3, N);
    u = zeros(3, N);
    A_vals = zeros(3,3,N);
    B_vals = zeros(3,3,N);
    L_vals = zeros(6,6,N);
    Q_vals = zeros(6,3,N);
    A_fro_vals = zeros(1,N);
    B_fro_vals = zeros(1,N);
    cost = 0;
    for i = 1:N
        if i == 1
            x(:,i) = xInit;
            A_vals(:, :, i) = A0;
            B_vals(:, :, i) = B0;

```

```

        L_vals(:,i) = eye(6);
        Q_vals(:,i) = [A0' B0'];
    else
        noise = normrnd(0, noiseStdDev, [3,1]);
        xk = x(:,i-1);
        uk = u(:,i-1);
        Ak = A_vals(:,i-1);
        Bk = B_vals(:,i-1);
        x(:,i) = A*xk + B*uk + noise;

        % Update L, Q, A, B vals for ith
        xbar = [xk' uk'];
        Lk = L_vals(:,i-1);
        Qk = Q_vals(:,i-1);
        Lknext = Lk - (1/(1+xbar'*Lk*xbar))*(Lk*xbar)*(Lk*xbar)';
        Qknext = xbar*x(:,i)' + Qk;
        LQnext = (Lknext*Qknext)';
        L_vals(:,i) = Lknext;
        Q_vals(:,i) = Qknext;
        A_vals(:,i) = LQnext(1:3);
        B_vals(:,i) = LQnext(4:6);
    end

    [K,~,~] = dlqr(A_vals(:,i),B_vals(:,i),Q,R);
    u(:,i) = -K*x(:,i);
    cost = cost + x(:,i)'*Q*x(:,i) + u(:,i)'*R*u(:,i);

    % Store frobenius norms
    A_fro_vals(i) = norm(A_vals(:,i)-A,'fro');
    B_fro_vals(i) = norm(B_vals(:,i)-B,'fro');
end

fprintf('Cost for N = %d is %.2f \n',N, cost)

figure
hold on
plot([1:N], x(1,:), 'r')
plot([1:N], x(2,:), 'g')
plot([1:N], x(3,:), 'b')
titl = sprintf('State vs. time for N = %d',N);
legend({'$x_1$', '$x_2$', '$x_3$'}, 'Interpreter', 'latex', 'FontSize', 20, 'Location', 'east');
xlabel('Time', 'Interpreter', 'latex', 'FontSize', 20);
ylabel('State', 'Interpreter', 'latex', 'FontSize', 20);
title(titl, 'Interpreter', 'latex', 'FontSize', 20);
grid on

figure
hold on
plot([1:N], u(1,:), 'r')
plot([1:N], u(2,:), 'g')
plot([1:N], u(3,:), 'b')
titl = sprintf('Control vs. time for N = %d',N);
legend({'$u_1$', '$u_2$', '$u_3$'}, 'Interpreter', 'latex', 'FontSize', 20, 'Location', 'east');
xlabel('Time', 'Interpreter', 'latex', 'FontSize', 20);
ylabel('Control', 'Interpreter', 'latex', 'FontSize', 20);
title(titl, 'Interpreter', 'latex', 'FontSize', 20);
grid on

figure
hold on
plot([1:N], A_fro_vals, 'r')
plot([1:N], B_fro_vals, 'b')
titl = sprintf('Frobenius norms vs. time for N = %d',N);
legend({'$||\hat{A}_k-A||$', '$||\hat{B}_k-B||$'}, 'Interpreter', 'latex', 'FontSize', 20, 'Location', 'east');
xlabel('Time', 'Interpreter', 'latex', 'FontSize', 20);
ylabel('Frobenius norm', 'Interpreter', 'latex', 'FontSize', 20);
title(titl, 'Interpreter', 'latex', 'FontSize', 18);
grid on
end






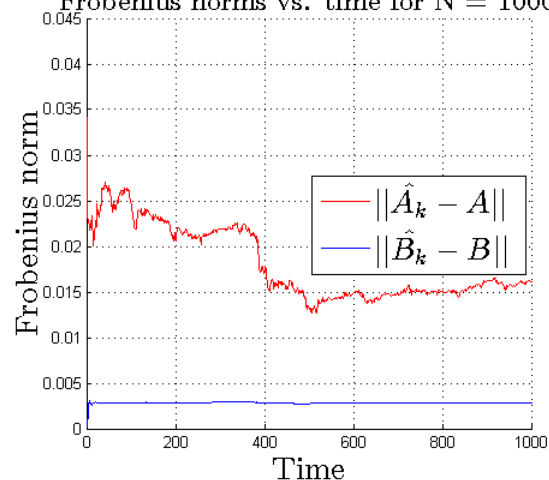
```

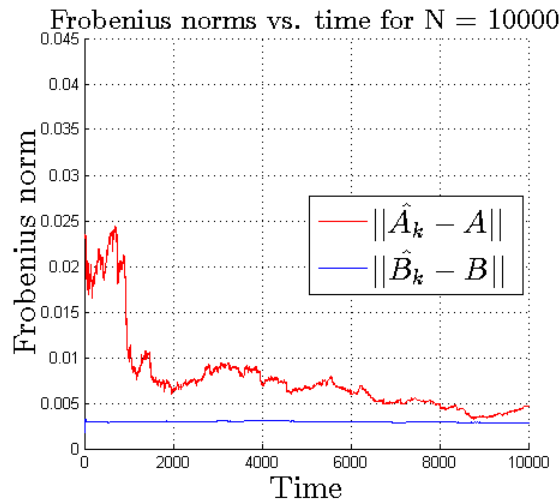
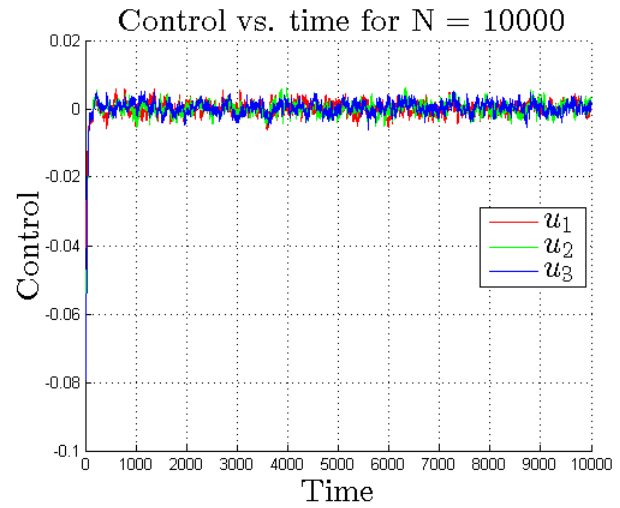
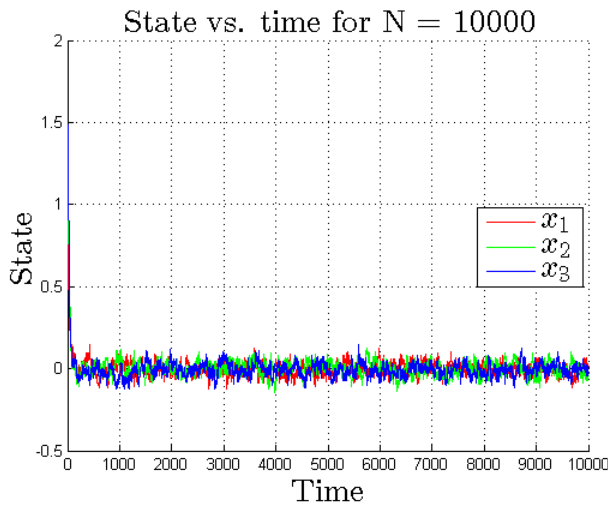
```

Cost for N = 100 is 478.99
Cost for N = 1000 is 499.20
Cost for N = 10000 is 601.38

```

After introducing a certainty-equivalent adaptive LQR controller, we see that in just a few timesteps, the control converges to 0 and the state is brought to 0. The cost is much lower than with the naive LQR controller. However, the system identification is not perfect because the Frobenius norm of the difference between the estimated parameter matrices and the true parameter matrices does not converge to 0.

State vs. time for $N = 100$ Control vs. time for $N = 100$ Frobenius norms vs. time for $N = 100$ State vs. time for $N = 1000$ Control vs. time for $N = 1000$ Frobenius norms vs. time for $N = 1000$ 



Part c- adaptive LQR controller with white noise

```

N_vals = [100 1000 10000];
white_noise_stddev_vals = [0.00001 0.001 0.01 1];
for white_noise_stddev = white_noise_stddev_vals
    for N = N_vals
        x = zeros(3, N);
        u = zeros(3, N);
        A_vals = zeros(3,3,N);
        B_vals = zeros(3,3,N);
        L_vals = zeros(6,6,N);
        Q_vals = zeros(6,6,N);
        A_fro_vals = zeros(1,N);
        B_fro_vals = zeros(1,N);
        cost = 0;
        for i = 1:N
            if i == 1
                x(:,i) = xInit;
                A_vals(:, :, i) = A0;
                B_vals(:, :, i) = B0;
                L_vals(:, :, i) = eye(6);
                Q_vals(:, :, i) = [A0' B0'];
            else
                noise = normrnd(0, noiseStdDev, [3,1]);
                xk = x(:,i-1);
                uk = u(:,i-1);
                Ak = A_vals(:, :, i-1);
                Bk = B_vals(:, :, i-1);
                x(:,i) = A*xk + B*uk + noise;

                % Update L, Q, A, B vals for ith
                xbar = [xk' uk']';
                Lk = L_vals(:, :, i-1);
                Qk = Q_vals(:, :, i-1);
                Lknext = Lk - (1/(1+xbar'*Lk*xbar))*(Lk*xbar)*(Lk*xbar)';
                Qknext = xbar*x(:,i)' + Qk;
                LQnext = (Lknext*Qknext)';
                L_vals(:, :, i) = Lknext;
                Q_vals(:, :, i) = Qknext;
                A_vals(:, :, i) = LQnext(:,1:3);
                B_vals(:, :, i) = LQnext(:,4:6);
            end
            white_noise = normrnd(0, white_noise_stddev, [3,1]);
            [K,~,~] = dlqr(A_vals(:, :, i), B_vals(:, :, i), Q, R);
            u(:,i) = -K*x(:,i) + white_noise;
            cost = cost + x(:,i)'*Q*x(:,i) + u(:,i)'*R*u(:,i);
        end
    end
end

```

```

        % Store frobenius norms
        A_fro_vals(i) = norm(A_vals(:,i)-A,'fro');
        B_fro_vals(i) = norm(B_vals(:,i)-B,'fro');
    end

fprintf('Cost for N = %d, white noise std dev = %.5f is %.2f \n',N, white_noise_stddev, cost)

figure
hold on
plot([1:N], x(1,:), 'r')
plot([1:N], x(2,:), 'g')
plot([1:N], x(3,:), 'b')
titl = sprintf('State vs. time for N = %d $\sigma = %.5f$',N, white_noise_stddev);
legend({'$x_1$', '$x_2$', '$x_3$'}, 'Interpreter', 'latex', 'FontSize', 20, 'Location', 'east');
xlabel('Time', 'Interpreter', 'latex', 'FontSize', 20);
ylabel('State', 'Interpreter', 'latex', 'FontSize', 20);
title(titl, 'Interpreter', 'latex', 'FontSize', 20);
grid on

figure
hold on
plot([1:N], u(1,:), 'r')
plot([1:N], u(2,:), 'g')
plot([1:N], u(3,:), 'b')
titl = sprintf('Control vs. time for N = %d $\sigma = %.5f$',N, white_noise_stddev);
legend({'$u_1$', '$u_2$', '$u_3$'}, 'Interpreter', 'latex', 'FontSize', 20, 'Location', 'east');
xlabel('Time', 'Interpreter', 'latex', 'FontSize', 20);
ylabel('Control', 'Interpreter', 'latex', 'FontSize', 20);
title(titl, 'Interpreter', 'latex', 'FontSize', 20);
grid on

figure
hold on
plot([1:N], A_fro_vals, 'r')
plot([1:N], B_fro_vals, 'b')
titl = sprintf('Frobenius norms vs. time for N = %d $\sigma = %.5f$',N, white_noise_stddev);
legend({'$\hat{A}_k-A$', '$\hat{B}_k-B$'}, 'Interpreter', 'latex', 'FontSize', 20, 'Location', 'east');
xlabel('Time', 'Interpreter', 'latex', 'FontSize', 20);
ylabel('Frobenius norm', 'Interpreter', 'latex', 'FontSize', 20);
title(titl, 'Interpreter', 'latex', 'FontSize', 18);
grid on
end
end


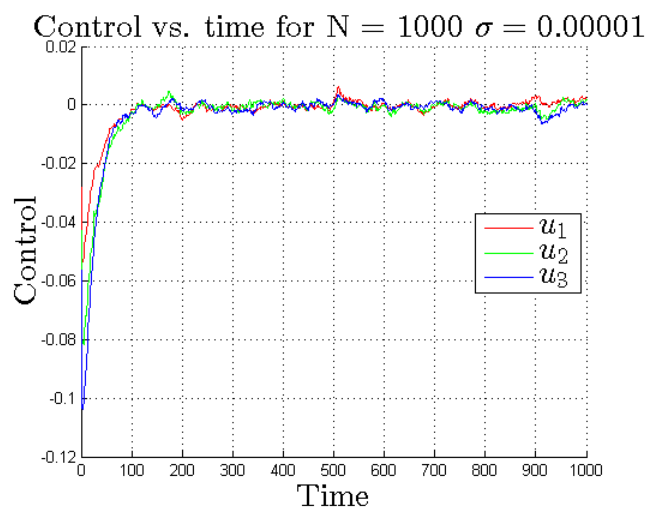




```

```

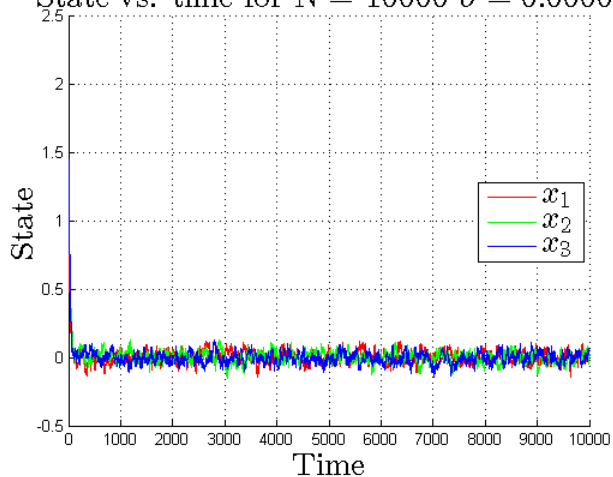
Cost for N = 100, white noise std dev = 0.00001 is 492.18
Cost for N = 1000, white noise std dev = 0.00001 is 482.81
Cost for N = 10000, white noise std dev = 0.00001 is 635.98
Cost for N = 100, white noise std dev = 0.00100 is 471.72
Cost for N = 1000, white noise std dev = 0.00100 is 495.86
Cost for N = 10000, white noise std dev = 0.00100 is 620.78
Cost for N = 100, white noise std dev = 0.01000 is 509.27
Cost for N = 1000, white noise std dev = 0.01000 is 777.46
Cost for N = 10000, white noise std dev = 0.01000 is 3776.37
Cost for N = 100, white noise std dev = 1.00000 is 310598.40
Cost for N = 1000, white noise std dev = 1.00000 is 3097063.33
Cost for N = 10000, white noise std dev = 1.00000 is 31337434.89

```


We notice a few interesting things after introducing white noise to the control. The first is that when the white noise standard deviation is very small, the performance is similar to the certainty equivalent LQR controller (just like we'd expect). As the white noise gets larger, the control input gets more jittery, and the cost increases (without blowing up), but the controller also does more exploration so the system identification gets better. At large standard deviation, with $\sigma = 1$, we see that the system identification becomes perfect, i.e. the Frobenius norm of the difference between the estimated parameter matrix and the true matrix converges to 0 (see last plot).



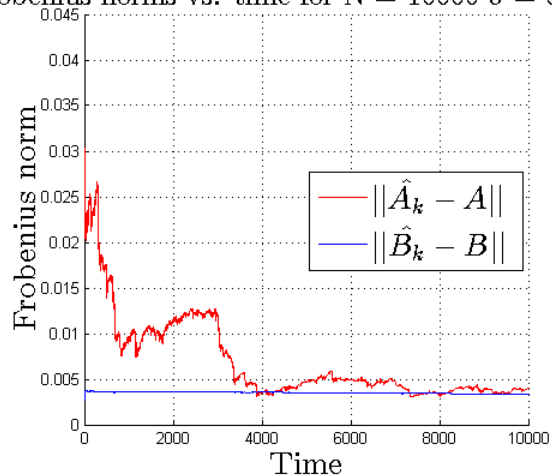
State vs. time for $N = 10000$ $\sigma = 0.00001$




Control vs. time for $N = 10000$ $\sigma = 0.00001$



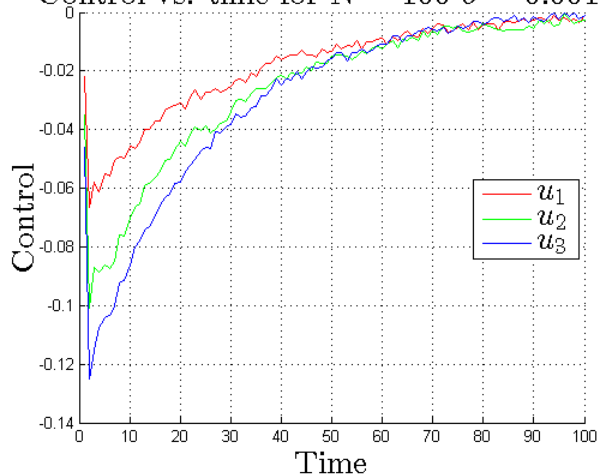
Frobenius norms vs. time for $N = 10000$ $\sigma = 0.00$



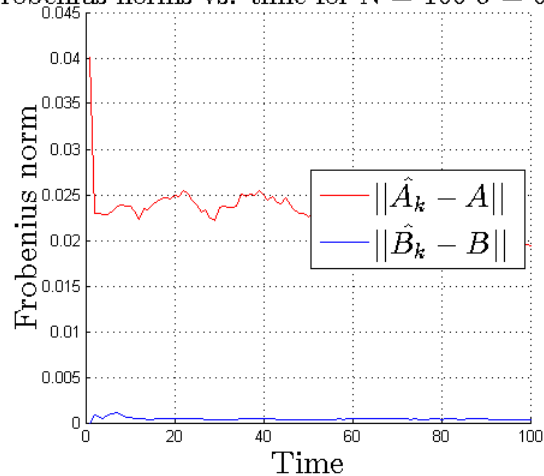
State vs. time for $N = 100$ $\sigma = 0.00100$




Control vs. time for $N = 100$ $\sigma = 0.00100$




Frobenius norms vs. time for $N = 100$ $\sigma = 0.001$




State vs. time for $N = 1000$ $\sigma = 0.00100$




Control vs. time for $N = 1000$ $\sigma = 0.00100$




Frobenius norms vs. time for $N = 1000$ $\sigma = 0.00$




State vs. time for $N = 10000$ $\sigma = 0.00100$









Control vs. time for $N = 10000$ $\sigma = 0.00100$




Frobenius norms vs. time for $N = 10000$ $\sigma = 0.00$







State vs. time for $N = 10000$ $\sigma = 0.01000$




Control vs. time for $N = 10000$ $\sigma = 0.01000$




Frobenius norms vs. time for $N = 10000$ $\sigma = 0.01$



State vs. time for $N = 100$ $\sigma = 1.00000$



Control vs. time for $N = 100$ $\sigma = 1.00000$



Frobenius norms vs. time for $N = 100$ $\sigma = 1.000$

